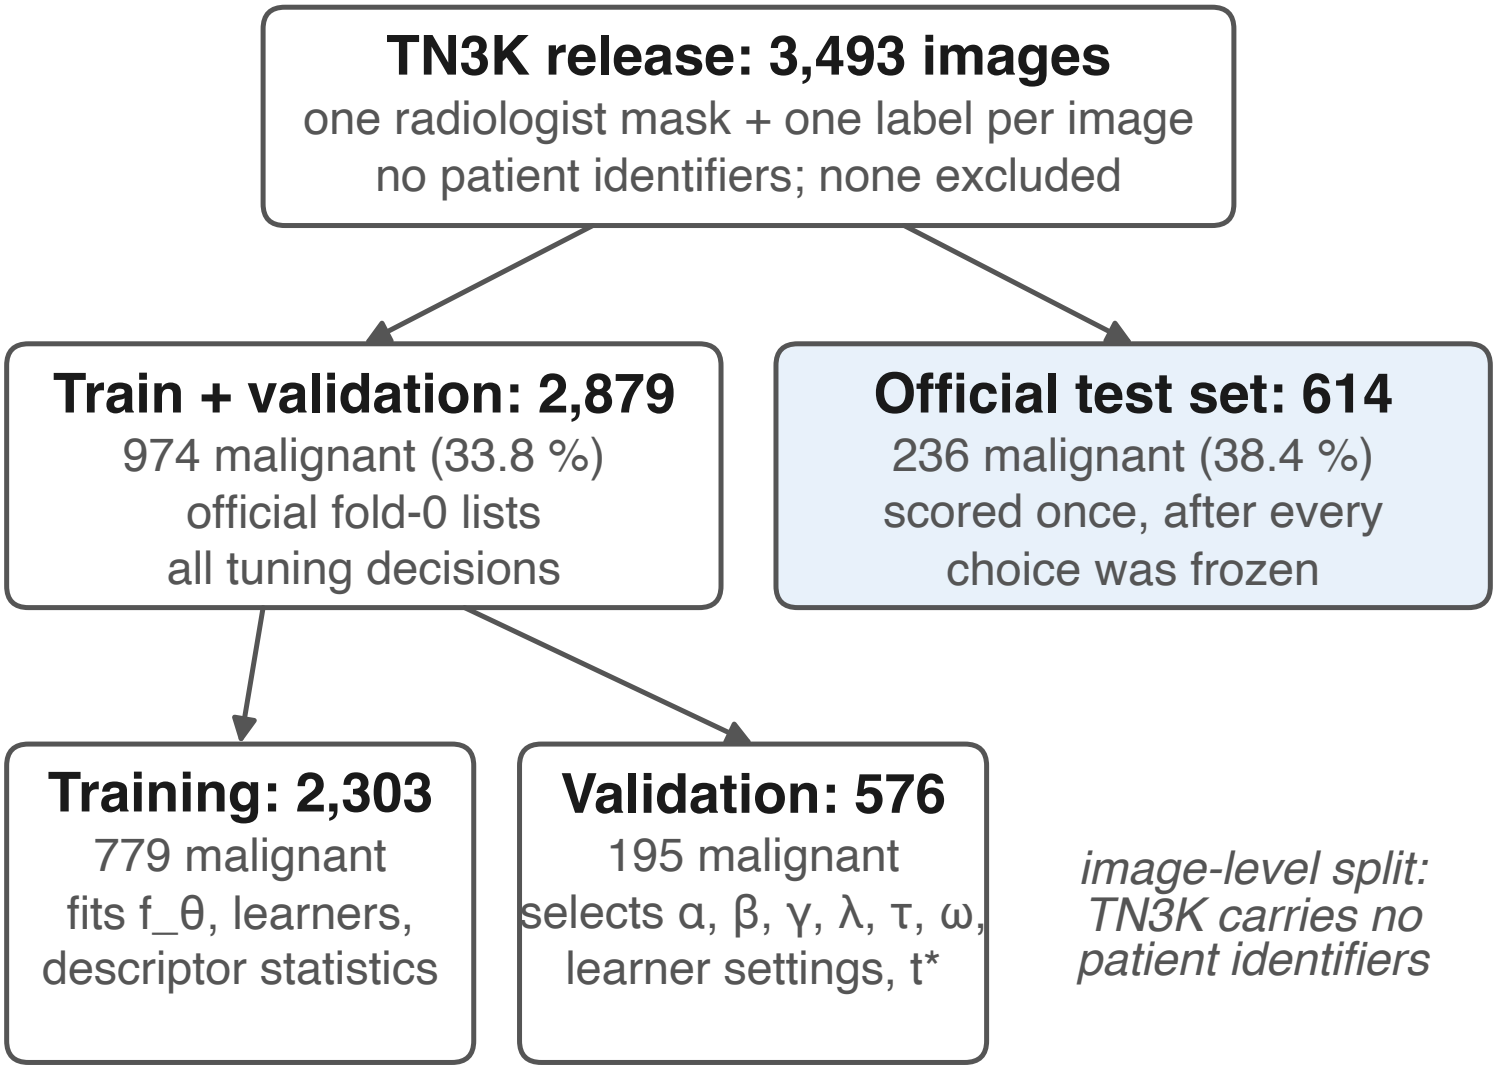
TN3K release: 3,493 images
one radiologist mask + one label per image
no patient identifiers; none excluded
Train + validation: 2,879
974 malignant (33.8 %)
official fold-0 lists
all tuning decisions
Official test set: 614
236 malignant (38.4 %)
scored once, after every
choice was frozen
Training: 2,303
779 malignant
fits f_θ, learners,
descriptor statistics
Validation: 576
195 malignant
selects α, β, γ, λ, τ, ω,
learner settings, t*
image-level split:
TN3K carries no
patient identifiers